

Data Wrangling

Estimated time needed: **30** minutes

Objectives

After completing this lab you will be able to:

- Handle missing values
- Correct data formatting
- Standardize and normalize data

You use data wrangling to convert data from an initial format to a format that may be better for analysis.

```
#install specific version of libraries used in lab  
#! mamba install pandas==1.3.3  
#! mamba install numpy=1.21.2  
  
import pandas as pd  
import matplotlib.pyplot as plt
```

The functions below will download the dataset into your browser:

```
from pyodide.http import pyfetch  
  
async def download(url, filename):  
    response = await pyfetch(url)  
    if response.status == 200:  
        with open(filename, "wb") as f:  
            f.write(await response.bytes())
```

First, assign the URL of the data set to "filepath".

```
file_path="https://cf-courses-data.s3.us.cloud-object-  
storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DA0101EN-  
SkillsNetwork/labs/Data%20files/auto.csv"
```

To obtain the dataset, utilize the download() function as defined above:

```
await download(file_path, "usedcars.csv")  
file_name="usedcars.csv"
```

Then, create a Python list headers containing name of headers.

```
headers = ["symboling", "normalized-losses", "make", "fuel-
type", "aspiration", "num-of-doors", "body-style",
          "drive-wheels", "engine-location", "wheel-base",
"length", "width", "height", "curb-weight", "engine-type",
          "num-of-cylinders", "engine-size", "fuel-
system", "bore", "stroke", "compression-ratio", "horsepower",
          "peak-rpm", "city-mpg", "highway-mpg", "price"]
```

Use the Pandas method `read_csv()` to load the data from the web address. Set the parameter "names" equal to the Python list "headers".

```
df = pd.read_csv(file_name, names = headers, on_bad_lines='skip')
```

Note: This version of the lab is working on JupyterLite, which requires the dataset to be downloaded to the interface. While working on the downloaded version of this notebook on their local machines (Jupyter Anaconda), the learners can simply **skip the steps above**, and simply use the URL directly in the `pandas.read_csv()` function. You can uncomment and run the statements in the cell below.

```
#filepath = "https://cf-courses-data.s3.us.cloud-object-
storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DA0101EN-
SkillsNetwork/labs/Data%20files/auto.csv"
#df = pd.read_csv(filepath, header=headers)    # Utilize the same
header list defined above
```

Use the method `head()` to display the first five rows of the dataframe.

```
# To see what the data set looks like, we'll use the head() method.
df.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-
	of-doors	\				
0	3.0	?	alfa-romero	gas	std	
two						
1	3.0	?	alfa-romero	gas	std	
two						
2	1.0	?	alfa-romero	gas	std	
two						
3	2.0	164	audi	gas	std	
four						
4	2.0	164	audi	gas	std	
four						

	body-style	drive-wheels	engine-location	wheel-base	...	engine-
	size	\				
0	convertible	rwd	front	88.6	...	
130.0						
1	convertible	rwd	front	88.6	...	
130.0						

```

2    hatchback      rwd      front      94.5 ...
152.0
3         sedan      fwd      front      99.8 ...
109.0
4         sedan      4wd      front      99.4 ...
136.0

    fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm
city-mpg \
0      mpfi  3.47    2.68              9.0         111      5000
21.0
1      mpfi  3.47    2.68              9.0         111      5000
21.0
2      mpfi  2.68    3.47              9.0         154      5000
19.0
3      mpfi  3.19    3.40             10.0         102      5500
24.0
4      mpfi  3.19    3.40              8.0         115      5500
18.0

    highway-mpg  price
0         27.0   13495
1         27.0   16500
2         26.0   16500
3         30.0   13950
4         22.0   17450

[5 rows x 26 columns]

```

As you can see, several question marks appeared in the data frame; those missing values may hinder further analysis. So, how do we identify all those missing values and deal with them?

How to work with missing data?

Steps for working with missing data: Identify missing data Deal with missing data Correct data format

Identify and handle missing values

Identify missing values

In the car data set, missing data comes with the question mark "?". We replace "?" with NaN (Not a Number), Python's default missing value marker for reasons of computational speed and convenience. Use the function: `.replace(A, B, inplace = True)` to replace A by B.

```

import numpy as np

# replace "?" to NaN

```

```
df.replace("?", np.nan, inplace = True)
df.head(5)
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-
of-doors \						
0	3.0	NaN	alfa-romero	gas	std	
two						
1	3.0	NaN	alfa-romero	gas	std	
two						
2	1.0	NaN	alfa-romero	gas	std	
two						
3	2.0	164	audi	gas	std	
four						
4	2.0	164	audi	gas	std	
four						

	body-style	drive-wheels	engine-location	wheel-base	...	engine-
size \						
0	convertible	rwd	front	88.6	...	
130.0						
1	convertible	rwd	front	88.6	...	
130.0						
2	hatchback	rwd	front	94.5	...	
152.0						
3	sedan	fwd	front	99.8	...	
109.0						
4	sedan	4wd	front	99.4	...	
136.0						

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm
city-mpg \						
0	mpfi	3.47	2.68	9.0	111	5000
21.0						
1	mpfi	3.47	2.68	9.0	111	5000
21.0						
2	mpfi	2.68	3.47	9.0	154	5000
19.0						
3	mpfi	3.19	3.40	10.0	102	5500
24.0						
4	mpfi	3.19	3.40	8.0	115	5500
18.0						

	highway-mpg	price
0	27.0	13495
1	27.0	16500
2	26.0	16500
3	30.0	13950
4	22.0	17450

[5 rows x 26 columns]

The missing values are converted by default. Use the following functions to identify these missing values. You can use two methods to detect missing data: `.isnull()` `.notnull()` The output is a boolean value indicating whether the value that is passed into the argument is in fact missing data.

```
missing_data = df.isnull()
missing_data.head(5)
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors
0	False	True	False	False	False	False
1	False	True	False	False	False	False
2	False	True	False	False	False	False
3	False	False	False	False	False	False
4	False	False	False	False	False	False

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size
0	False	False	False	False	...	False
1	False	False	False	False	...	False
2	False	False	False	False	...	False
3	False	False	False	False	...	False
4	False	False	False	False	...	False

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm
0	False	False	False	False	False	False
1	False	False	False	False	False	False
2	False	False	False	False	False	False
3	False	False	False	False	False	False
4	False	False	False	False	False	False

	city-mpg	highway-mpg	price
0	False	False	False
1	False	False	False

2	False	False	False
3	False	False	False
4	False	False	False

[5 rows x 26 columns]

"True" means the value is a missing value while "False" means the value is not a missing value.

```
for column in missing_data.columns.values.tolist():
    print(column)
    print(missing_data[column].value_counts())
    print("")
```

symboling

False 206

Name: symboling, dtype: int64

normalized-losses

False 165

True 41

Name: normalized-losses, dtype: int64

make

False 206

Name: make, dtype: int64

fuel-type

False 206

Name: fuel-type, dtype: int64

aspiration

False 206

Name: aspiration, dtype: int64

num-of-doors

False 204

True 2

Name: num-of-doors, dtype: int64

body-style

False 206

Name: body-style, dtype: int64

drive-wheels

False 206

Name: drive-wheels, dtype: int64

engine-location

False 206

Name: engine-location, dtype: int64

```
wheel-base
False      206
Name: wheel-base, dtype: int64

length
False      206
Name: length, dtype: int64

width
False      206
Name: width, dtype: int64

height
False      206
Name: height, dtype: int64

curb-weight
False      206
Name: curb-weight, dtype: int64

engine-type
False      206
Name: engine-type, dtype: int64

num-of-cylinders
False      206
Name: num-of-cylinders, dtype: int64

engine-size
False      206
Name: engine-size, dtype: int64

fuel-system
False      206
Name: fuel-system, dtype: int64

bore
False      202
True        4
Name: bore, dtype: int64

stroke
False      202
True        4
Name: stroke, dtype: int64

compression-ratio
False      205
True        1
```

```
Name: compression-ratio, dtype: int64
```

```
horsepower
```

```
False      203
```

```
True        3
```

```
Name: horsepower, dtype: int64
```

```
peak-rpm
```

```
False      203
```

```
True        3
```

```
Name: peak-rpm, dtype: int64
```

```
city-mpg
```

```
False      205
```

```
True        1
```

```
Name: city-mpg, dtype: int64
```

```
highway-mpg
```

```
False      205
```

```
True        1
```

```
Name: highway-mpg, dtype: int64
```

```
price
```

```
False      201
```

```
True        5
```

```
Name: price, dtype: int64
```

Based on the summary above, each column has 205 rows of data and seven of the columns containing missing data: "normalized-losses": 41 missing data "num-of-doors": 2 missing data "bore": 4 missing data "stroke": 4 missing data "horsepower": 2 missing data "peak-rpm": 2 missing data "price": 4 missing data

Deal with missing data

How should you deal with missing data?

You should only drop whole columns if most entries in the column are empty. In the data set, none of the columns are empty enough to drop entirely. You have some freedom in choosing which method to replace data; however, some methods may seem more reasonable than others. Apply each method to different columns:

Replace by mean: "normalized-losses": 41 missing data, replace them with mean "stroke": 4 missing data, replace them with mean "bore": 4 missing data, replace them with mean "horsepower": 2 missing data, replace them with mean "peak-rpm": 2 missing data, replace them with mean

Replace by frequency: "num-of-doors": 2 missing data, replace them with "four". Reason: 84% sedans are four doors. Since four doors is most frequent, it is most likely to occur

Drop the whole row: "price": 4 missing data, simply delete the whole row Reason: You want to predict price. You cannot use any data entry without price data for prediction; therefore any row now without price data is not useful to you.

```
avg_norm_loss = df["normalized-losses"].astype("float").mean(axis=0)
print("Average of normalized-losses:", avg_norm_loss)
```

```
Average of normalized-losses: 121.26611111111111
```

```
df["normalized-losses"].replace(np.nan, avg_norm_loss, inplace=True)
```

```
avg_bore=df['bore'].astype('float').mean(axis=0)
print("Average of bore:", avg_bore)
```

```
Average of bore: 3.313267326732673
```

```
df["bore"].replace(np.nan, avg_bore, inplace=True)
```

```
# Write your code below and press Shift+Enter to execute
#Calculate the mean vaule for "stroke" column
```

```
avg_stroke = df["stroke"].astype("float").mean(axis = 0)
print("Average of stroke:", avg_stroke)
```

```
# replace NaN by mean value in "stroke" column
```

```
df["stroke"].replace(np.nan, avg_stroke, inplace = True)
```

```
Average of stroke: 3.256903553299492
```

```
avg_horsepower = df['horsepower'].astype('float').mean(axis=0)
print("Average horsepower:", avg_horsepower)
```

```
Average horsepower: 104.25615763546799
```

```
df['horsepower'].replace(np.nan, avg_horsepower, inplace=True)
```

```
avg_peakrpm=df['peak-rpm'].astype('float').mean(axis=0)
print("Average peak rpm:", avg_peakrpm)
```

```
Average peak rpm: 5125.369458128079
```

```
df['peak-rpm'].replace(np.nan, avg_peakrpm, inplace=True)
```

To see which values are present in a particular column, we can use the ".value_counts()" method:

```
df['num-of-doors'].value_counts()
```

```
four      115
two        89
```

```
Name: num-of-doors, dtype: int64
```

You can see that four doors is the most common type. We can also use the ".idxmax()" method to calculate the most common type automatically:

```
df['num-of-doors'].value_counts().idxmax()
'four'
```

The replacement procedure is very similar to what you have seen previously:

```
#replace the missing 'num-of-doors' values by the most frequent
df["num-of-doors"].replace(np.nan, "four", inplace=True)
```

Finally, drop all rows that do not have price data:

```
# simply drop whole row with NaN in "price" column
df.dropna(subset=["price"], axis=0, inplace=True)

# reset index, because we dropped two rows
df.reset_index(drop=True, inplace=True)

df.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors \
0	3.0	121.266111	alfa-romero	gas	std	two
1	3.0	121.266111	alfa-romero	gas	std	two
2	1.0	121.266111	alfa-romero	gas	std	two
3	2.0	164	audi	gas	std	four
4	2.0	164	audi	gas	std	four

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size \
0	convertible	rwd	front	88.6	...	130.0
1	convertible	rwd	front	88.6	...	130.0
2	hatchback	rwd	front	94.5	...	152.0
3	sedan	fwd	front	99.8	...	109.0
4	sedan	4wd	front	99.4	...	136.0

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg \
0	mpfi	3.47	2.68	9.0	111	5000	21.0
1	mpfi	3.47	2.68	9.0	111	5000	

21.0						
2	mpfi	2.68	3.47	9.0	154	5000
19.0						
3	mpfi	3.19	3.40	10.0	102	5500
24.0						
4	mpfi	3.19	3.40	8.0	115	5500
18.0						

	highway-mpg	price
0	27.0	13495
1	27.0	16500
2	26.0	16500
3	30.0	13950
4	22.0	17450

[5 rows x 26 columns]

Good! Now, you have a data set with no missing values.

Correct data format

We are almost there! The last step in data cleaning is checking and making sure that all data is in the correct format (int, float, text or other).

In Pandas, you use: `.dtype()` to check the data type `.astype()` to change the data type

```
df.dtypes
```

symboling	float64
normalized-losses	object
make	object
fuel-type	object
aspiration	object
num-of-doors	object
body-style	object
drive-wheels	object
engine-location	object
wheel-base	float64
length	float64
width	float64
height	float64
curb-weight	int64
engine-type	object
num-of-cylinders	object
engine-size	float64
fuel-system	object
bore	object
stroke	object
compression-ratio	float64
horsepower	object

```

peak-rpm          object
city-mpg          float64
highway-mpg       float64
price             object
dtype: object

df[["bore", "stroke"]] = df[["bore", "stroke"]].astype("float")
df[["normalized-losses"]] = df[["normalized-losses"]].astype("int")
df[["price"]] = df[["price"]].astype("float")
df[["peak-rpm"]] = df[["peak-rpm"]].astype("float")

df.dtypes

symboling         float64
normalized-losses  int32
make              object
fuel-type         object
aspiration        object
num-of-doors      object
body-style        object
drive-wheels      object
engine-location   object
wheel-base       float64
length            float64
width             float64
height            float64
curb-weight       int64
engine-type       object
num-of-cylinders  object
engine-size       float64
fuel-system       object
bore              float64
stroke            float64
compression-ratio float64
horsepower        object
peak-rpm          float64
city-mpg          float64
highway-mpg       float64
price             float64
dtype: object

```

Wonderful!

Now you finally obtained the cleansed data set with no missing values and with all data in its proper format.

Data Standardization

What is standardization? Standardization is the process of transforming data into a common format, allowing the researcher to make the meaningful comparison.

Example Transform mpg to L/100km: In your data set, the fuel consumption columns "city-mpg" and "highway-mpg" are represented by mpg (miles per gallon) unit. Assume you are developing an application in a country that accepts the fuel consumption with L/100km standard. You will need to apply data transformation to transform mpg into L/100km.

Use this formula for unit conversion: $L/100km = 235 / mpg$ You can do many mathematical operations directly using Pandas.

```
df.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	\
0	3.0	121	alfa-romero	gas	std	
1	3.0	121	alfa-romero	gas	std	
2	1.0	121	alfa-romero	gas	std	
3	2.0	164	audi	gas	std	
4	2.0	164	audi	gas	std	

	num-of-doors	body-style	drive-wheels	engine-location	wheel-
base ... \					
0	two	convertible	rwd	front	
88.6 ...					
1	two	convertible	rwd	front	
88.6 ...					
2	two	hatchback	rwd	front	
94.5 ...					
3	four	sedan	fwd	front	
99.8 ...					
4	four	sedan	4wd	front	
99.4 ...					

	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower
\						
0	130.0	mpfi	3.47	2.68	9.0	111
1	130.0	mpfi	3.47	2.68	9.0	111
2	152.0	mpfi	2.68	3.47	9.0	154
3	109.0	mpfi	3.19	3.40	10.0	102
4	136.0	mpfi	3.19	3.40	8.0	115

	peak-rpm	city-mpg	highway-mpg	price
0	5000.0	21.0	27.0	13495.0
1	5000.0	21.0	27.0	16500.0
2	5000.0	19.0	26.0	16500.0
3	5500.0	24.0	30.0	13950.0
4	5500.0	18.0	22.0	17450.0

```
[5 rows x 26 columns]
```

```
# Convert mpg to L/100km by mathematical operation (235 divided by mpg)
```

```
df['city-L/100km'] = 235/df["city-mpg"]
```

```
# check your transformed data
```

```
df.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	\
0	3.0	121	alfa-romero	gas	std	
1	3.0	121	alfa-romero	gas	std	
2	1.0	121	alfa-romero	gas	std	
3	2.0	164	audi	gas	std	
4	2.0	164	audi	gas	std	

	num-of-doors	body-style	drive-wheels	engine-location	wheel-
base ... \					
0	two	convertible	rwd	front	
88.6 ...					
1	two	convertible	rwd	front	
88.6 ...					
2	two	hatchback	rwd	front	
94.5 ...					
3	four	sedan	fwd	front	
99.8 ...					
4	four	sedan	4wd	front	
99.4 ...					

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm
city-mpg \						
0	mpfi	3.47	2.68	9.0	111	5000.0
21.0						
1	mpfi	3.47	2.68	9.0	111	5000.0
21.0						
2	mpfi	2.68	3.47	9.0	154	5000.0
19.0						
3	mpfi	3.19	3.40	10.0	102	5500.0
24.0						
4	mpfi	3.19	3.40	8.0	115	5500.0
18.0						

	highway-mpg	price	city-L/100km
0	27.0	13495.0	11.190476
1	27.0	16500.0	11.190476
2	26.0	16500.0	12.368421
3	30.0	13950.0	9.791667
4	22.0	17450.0	13.055556

```
[5 rows x 27 columns]
```

```

# Write your code below and press Shift+Enter to execute
# transform mpg to L/100km by mathematical operation (235 divided by mpg)
df["highway-mpg"] = 235/df["highway-mpg"]

# rename column name from "highway-mpg" to "highway-L/100km"
df.rename(columns={"highway-mpg":'highway-L/100km'}, inplace=True)

# check your transformed data
df.head()

```

	symboling	normalized-losses	make	aspiration	num-of-doors
0	3.0	121	alfa-romero	std	two
1	3.0	121	alfa-romero	std	two
2	1.0	121	alfa-romero	std	two
3	2.0	164	audi	std	four
4	2.0	164	audi	std	four

	body-style	drive-wheels	engine-location	wheel-base	length	...
0	convertible	rwd	front	88.6	0.811148	...
1	convertible	rwd	front	88.6	0.811148	...
2	hatchback	rwd	front	94.5	0.822681	...
3	sedan	fwd	front	99.8	0.848630	...
4	sedan	4wd	front	99.4	0.848630	...

	compression-ratio	horsepower	peak-rpm	city-mpg	highway-mpg
0	9.0	111	5000.0	21.0	8.703704
1	9.0	111	5000.0	21.0	8.703704
2	9.0	154	5000.0	19.0	9.038462
3	10.0	102	5500.0	24.0	7.833333
4	8.0	115	5500.0	18.0	10.681818

city-L/100km	horsepower-binned	fuel-type-diesel	fuel-type-gas
--------------	-------------------	------------------	---------------

0	11.190476	Low	0	1
1	11.190476	Low	0	1
2	12.368421	Medium	0	1
3	9.791667	Low	0	1
4	13.055556	Low	0	1

[5 rows x 29 columns]

Data Normalization

Why normalization? Normalization is the process of transforming values of several variables into a similar range. Typical normalizations include scaling the variable so the variable average is 0 scaling the variable so the variance is 1 scaling the variable so the variable values range from 0 to 1

Example To demonstrate normalization, say you want to scale the columns "length", "width" and "height". Target: normalize those variables so their value ranges from 0 to 1 Approach: replace the original value by (original value)/(maximum value)

```
# replace (original value) by (original value)/(maximum value)
df['length'] = df['length']/df['length'].max()
df['width'] = df['width']/df['width'].max()

# Write your code below and press Shift+Enter to execute
df['height'] = df['height']/df['height'].max()

# show the scaled columns
df[["length", "width", "height"]].head()
```

	length	width	height
0	0.811148	0.890278	0.816054
1	0.811148	0.890278	0.816054
2	0.822681	0.909722	0.876254
3	0.848630	0.919444	0.908027
4	0.848630	0.922222	0.908027

Here you've normalized "length", "width" and "height" to fall in the range of [0,1].

Binning

Why binning? Binning is a process of transforming continuous numerical variables into discrete categorical 'bins' for grouped analysis.

Example: In your data set, "horsepower" is a real valued variable ranging from 48 to 288 and it has 59 unique values. What if you only care about the price difference between cars with high horsepower, medium horsepower, and little horsepower (3 types)? You can rearrange them into three 'bins' to simplify analysis.

Convert data to correct format:

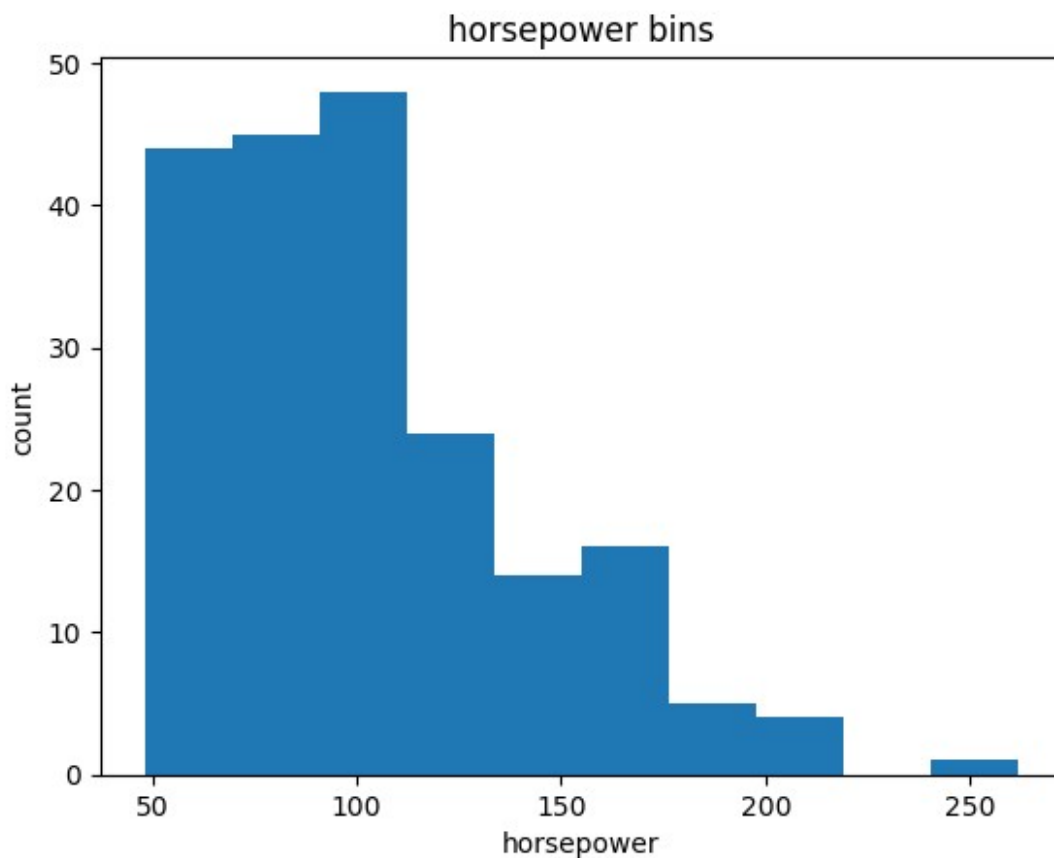

```
df["horsepower"]=df["horsepower"].astype(int, copy=True)
```

Plot the histogram of horsepower to see the distribution of horsepower.

```
%matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
plt.pyplot.hist(df["horsepower"])

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")

Text(0.5, 1.0, 'horsepower bins')
```



Build a bin array with a minimum value to a maximum value by using the bandwidth calculated above. The values will determine when one bin ends and another begins.

```
bins = np.linspace(min(df["horsepower"]), max(df["horsepower"]), 4)
bins
array([ 48.          , 119.33333333, 190.66666667, 262.          ])
```

Set group names:

```
group_names = ['Low', 'Medium', 'High']
```

Apply the function "cut" to determine what each value of `df['horsepower']` belongs to.

```
df['horsepower-binned'] = pd.cut(df['horsepower'], bins,  
labels=group_names, include_lowest=True )  
df[['horsepower', 'horsepower-binned']].head(20)
```

	horsepower	horsepower-binned
0	111	Low
1	111	Low
2	154	Medium
3	102	Low
4	115	Low
5	110	Low
6	110	Low
7	110	Low
8	140	Medium
9	101	Low
10	101	Low
11	121	Medium
12	121	Medium
13	121	Medium
14	182	Medium
15	182	Medium
16	182	Medium
17	48	Low
18	70	Low
19	70	Low

See the number of vehicles in each bin:

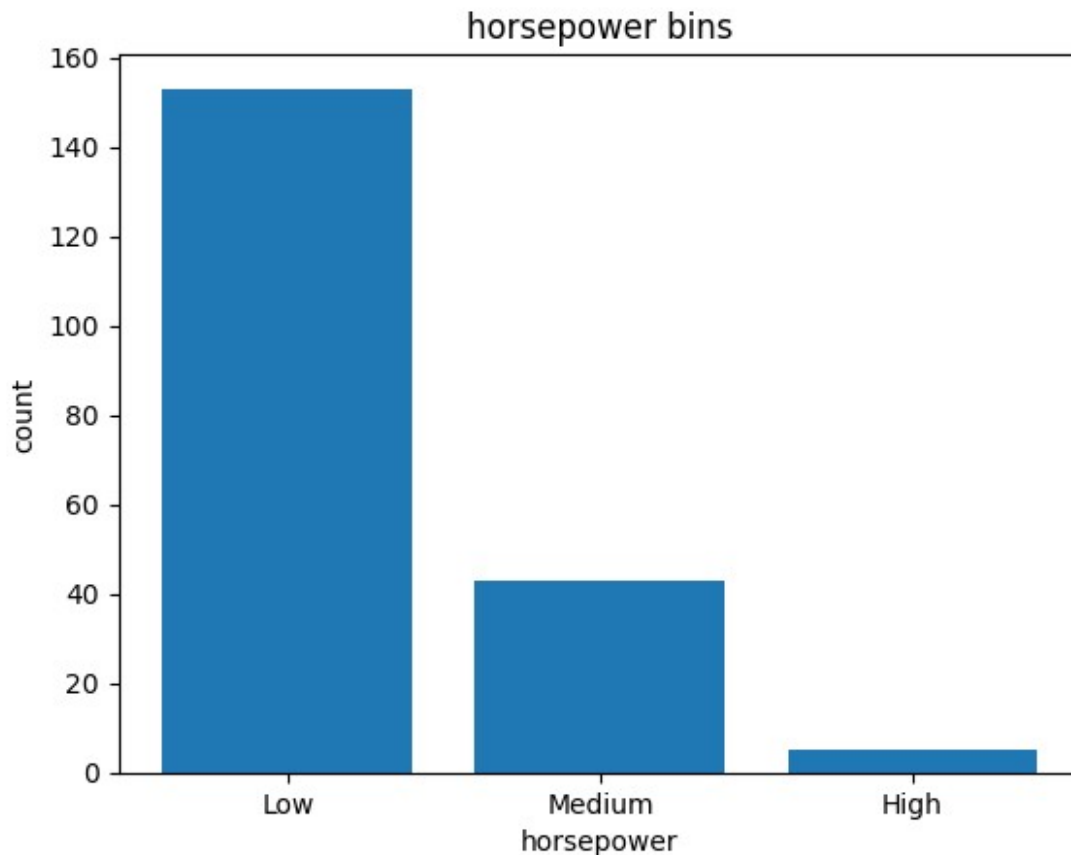
```
df["horsepower-binned"].value_counts()
```

```
Low      153  
Medium   43  
High      5  
Name: horsepower-binned, dtype: int64
```

Plot the distribution of each bin:

```
%matplotlib inline  
import matplotlib as plt  
from matplotlib import pyplot  
pyplot.bar(group_names, df["horsepower-binned"].value_counts())  
  
# set x/y labels and plot title
```

```
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
Text(0.5, 1.0, 'horsepower bins')
```

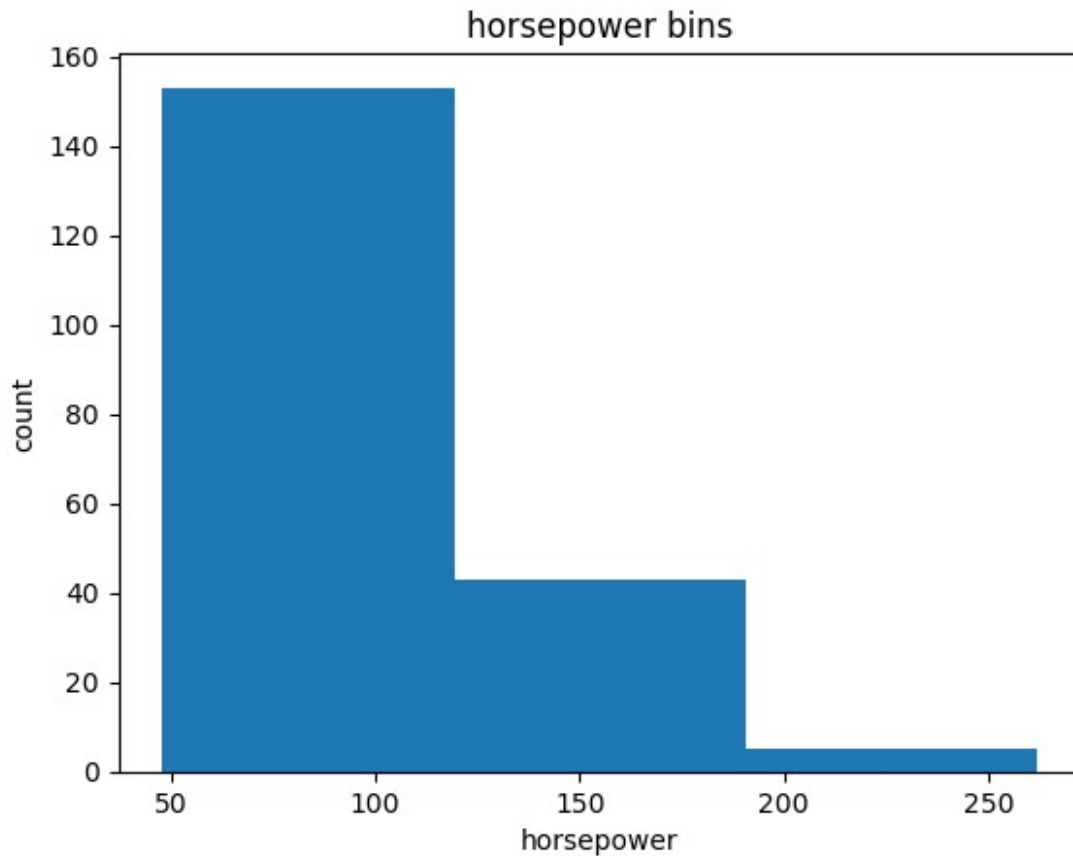


Normally, you use a histogram to visualize the distribution of bins we created above.

```
%matplotlib inline
import matplotlib as plt
from matplotlib import pyplot

# draw histogram of attribute "horsepower" with bins = 3
plt.pyplot.hist(df["horsepower"], bins = 3)

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
Text(0.5, 1.0, 'horsepower bins')
```



The plot above shows the binning result for the attribute "horsepower".

Indicator Variable

What is an indicator variable? An indicator variable (or dummy variable) is a numerical variable used to label categories. They are called 'dummies' because the numbers themselves don't have inherent meaning.

Why use indicator variables? You use indicator variables so you can use categorical variables for regression analysis in the later modules. Example The column "fuel-type" has two unique values: "gas" or "diesel". Regression doesn't understand words, only numbers. To use this attribute in regression analysis, you can convert "fuel-type" to indicator variables.

```
df.columns
Index(['symboling', 'normalized-losses', 'make', 'fuel-type',
      'aspiration',
      'num-of-doors', 'body-style', 'drive-wheels', 'engine-
location',
      'wheel-base', 'length', 'width', 'height', 'curb-weight',
      'engine-type',
      'num-of-cylinders', 'engine-size', 'fuel-system', 'bore',
      'stroke',
```

```
'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
'highway-mpg', 'price', 'city-L/100km', 'horsepower-binned'],
dtype='object')
```

Get the indicator variables and assign it to data frame "dummy_variable_1":

```
dummy_variable_1 = pd.get_dummies(df["fuel-type"])
dummy_variable_1.head()
```

	diesel	gas
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1

Change the column names for clarity:

```
dummy_variable_1.rename(columns={'gas': 'fuel-type-gas',
'diesel': 'fuel-type-diesel'}, inplace=True)
dummy_variable_1.head()
```

	fuel-type-diesel	fuel-type-gas
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1

In the data frame, column 'fuel-type' now has values for 'gas' and 'diesel' as 0s and 1s.

```
# merge data frame "df" and "dummy_variable_1"
df = pd.concat([df, dummy_variable_1], axis=1)
```

```
# drop original column "fuel-type" from "df"
df.drop("fuel-type", axis = 1, inplace=True)
```

```
df.head()
```

	symboling	normalized-losses	make	aspiration	num-of-doors
0	3.0	121	alfa-romero	std	two
1	3.0	121	alfa-romero	std	two
2	1.0	121	alfa-romero	std	two
3	2.0	164	audi	std	four
4	2.0	164	audi	std	four

	body-style	drive-wheels	engine-location	wheel-base	length	...
0	convertible	rwd	front	88.6	0.811148	...
1	convertible	rwd	front	88.6	0.811148	...
2	hatchback	rwd	front	94.5	0.822681	...
3	sedan	fwd	front	99.8	0.848630	...
4	sedan	4wd	front	99.4	0.848630	...

	compression-ratio	horsepower	peak-rpm	city-mpg	highway-mpg
0	9.0	111	5000.0	21.0	27.0
1	9.0	111	5000.0	21.0	27.0
2	9.0	154	5000.0	19.0	26.0
3	10.0	102	5500.0	24.0	30.0
4	8.0	115	5500.0	18.0	22.0

	city-L/100km	horsepower-binned	fuel-type-diesel	fuel-type-gas
0	11.190476	Low	0	1
1	11.190476	Low	0	1
2	12.368421	Medium	0	1
3	9.791667	Low	0	1
4	13.055556	Low	0	1

[5 rows x 29 columns]

The last two columns are now the indicator variable representation of the fuel-type variable. They're all 0s and 1s now.

```
# Write your code below and press Shift+Enter to execute
# get indicator variables of aspiration and assign it to data frame
"dummy_variable_2"
dummy_variable_2 = pd.get_dummies(df['aspiration'])

# change column names for clarity
dummy_variable_2.rename(columns={'std':'aspiration-std', 'turbo':
'aspiration-turbo'}, inplace=True)
```

```
# show first 5 instances of data frame "dummy_variable_1"
dummy_variable_2.head()
```

	aspiration-std	aspiration-turbo
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0

Question #5:

Merge the new dataframe to the original dataframe, then drop the column 'aspiration'.

```
# Write your code below and press Shift+Enter to execute
# merge the new dataframe to the original datafram
df = pd.concat([df, dummy_variable_2], axis=1)

# drop original column "aspiration" from "df"
df.drop('aspiration', axis = 1, inplace=True)
```

Save the new csv:

```
df.to_csv('clean_df.csv')
```

Thank you for completing this lab!

Author

Joseph Santarcangelo

Other Contributors

Mahdi Noorian PhD

Bahare Talayian

Eric Xiao

Steven Dong

Parizad

Hima Vasudevan

Fiorella Wenver

Yi Yao.

Abhishek Gagneja

Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2023-09-28	2.3	Abhishek Gagneja	Instructional Update
2020-10-30	2.2	Lakshmi	Changed URL of csv
2020-09-09	2.1	Lakshmi	Updated Indicator Variables section
2020-08-27	2.0	Lavanya	Moved lab to course repo in GitLab

© IBM Corporation 2023. All rights reserved.