

# Hands-on lab: Exploratory Data Analysis - Laptops Pricing dataset

Estimated time needed: **45** minutes

In this lab, you will use the skills acquired throughout the module, to explore the effect of different features on the price of laptops.

## Objectives

After completing this lab you will be able to:

- Visualize individual feature patterns
- Run descriptive statistical analysis on the dataset
- Use groups and pivot tables to find the effect of categorical variables on price
- Use Pearson Correlation to measure the interdependence between variables

## Setup

For this lab, we will be using the following libraries:

- `skillsnetwork` for downloading the data
- `pandas` for managing the data.
- `numpy` for mathematical operations.
- `scipy` for statistical operations.
- `seaborn` for visualizing the data.
- `matplotlib` for additional plotting tools.

## Install Required Libraries

You can install the required libraries by simply running the `pip install` command with a `%` sign before it. For this environment, `seaborn` library requires installation.

```
import piplite
await piplite.install('seaborn')
```

## Importing Required Libraries

*We recommend you import all required libraries in one place (here):*

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
%matplotlib inline
```

## Import the dataset

You should download the modified version of the data set from the last module. Run the following code block to download the CSV file to this environment.

The functions below will download the dataset into your browser:

```
from pyodide.http import pyfetch

async def download(url, filename):
    response = await pyfetch(url)
    if response.status == 200:
        with open(filename, "wb") as f:
            f.write(await response.bytes())

filepath="https://cf-courses-data.s3.us.cloud-object-
storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DA0101EN-Coursera/
laptop_pricing_dataset_mod2.csv"

await download(filepath, "laptops.csv")
file_name="laptops.csv"
```

Import the file to a pandas dataframe.

```
df = pd.read_csv(file_name, header=0)
```

Note: This version of the lab is working on JupyterLite, which requires the dataset to be downloaded to the interface. While working on the downloaded version of this notebook on their local machines, the learners can simply **skip the steps above**, and simply use the URL directly in the `pandas.read_csv()` function. You can uncomment and run the statements in the cell below.

```
#filepath="https://cf-courses-data.s3.us.cloud-object-
storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DA0101EN-Coursera/
laptop_pricing_dataset_mod2.csv"
#df = pd.read_csv(filepath, header=None)
```

Print the first 5 entries of the dataset to confirm loading.

```
df.head(5)
```

	Unnamed: 0.1	Unnamed: 0	Manufacturer	Category	GPU	OS	CPU_core
0	0	0	Acer	4	2	1	5
1	1	1	Dell	3	1	1	3
2	2	2	Dell	3	1	1	7
3	3	3	Dell	4	2	1	5
4	4	4	HP	4	2	1	7

	Screen_Size_inch	CPU_frequency	RAM_GB	Storage_GB_SSD
0	14.0	0.551724	8	256
1	15.6	0.689655	4	256
2	15.6	0.931034	8	256
3	13.3	0.551724	8	128
4	15.6	0.620690	8	256

	Price	Price-binned	Screen-Full_HD	Screen-IPS_panel
0	978	Low	0	1
1	634	Low	1	0
2	946	Low	1	0
3	1244	Low	0	1
4	837	Low	1	0

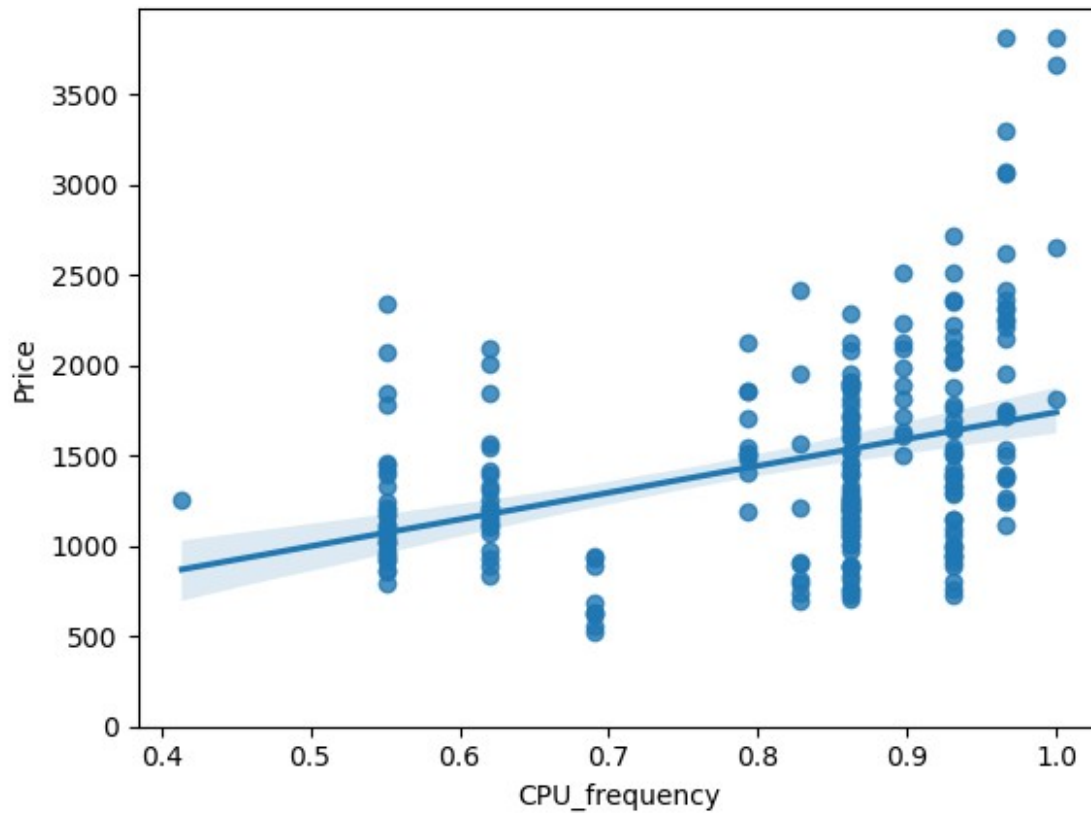
## Task 1 - Visualize individual feature patterns

### Continuous valued features

Generate regression plots for each of the parameters "CPU\_frequency", "Screen\_Size\_inch" and "Weight\_pounds" against "Price". Also, print the value of correlation of each feature with "Price".


```
# Write your code below and press Shift+Enter to execute
# CPU_frequency plot
sns.regplot(x="CPU_frequency", y="Price", data=df)
plt.ylim(0,)

(0.0, 3974.15)
```



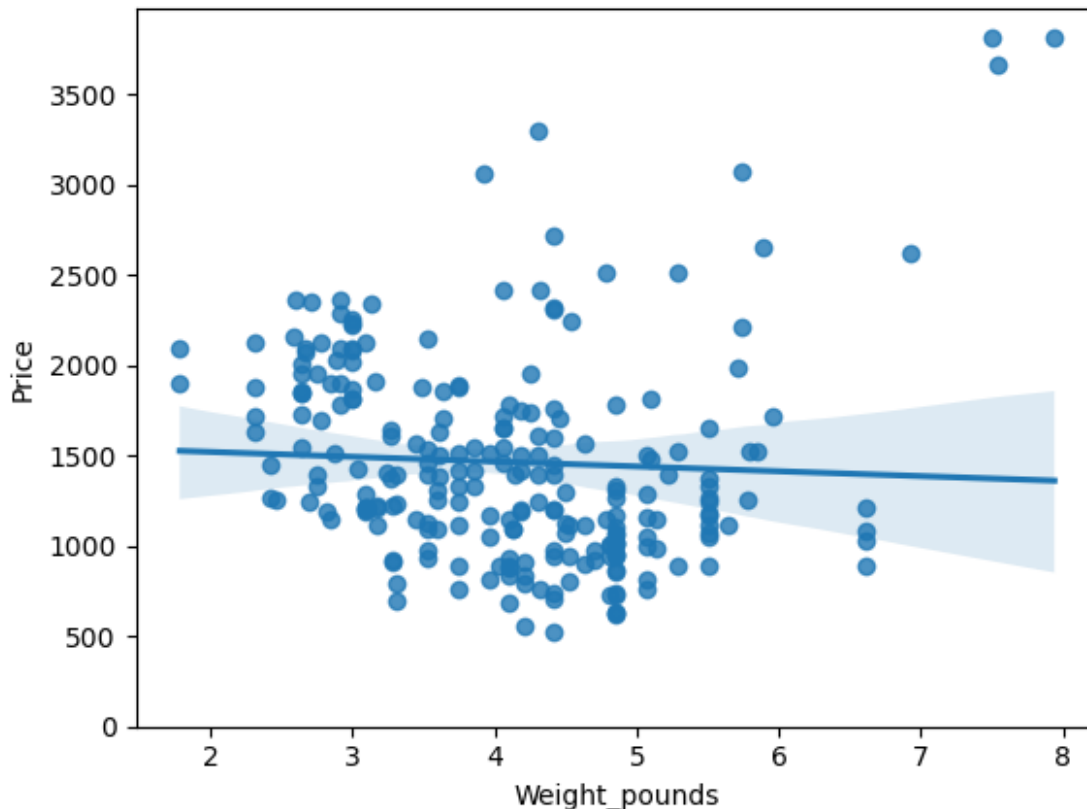
```
# Write your code below and press Shift+Enter to execute
# Screen_Size_inch plot
sns.regplot(x="Screen_Size_inch", y="Price", data=df)
plt.ylim(0,)

(0.0, 3974.15)
```



```
# Write your code below and press Shift+Enter to execute
# Weight_pounds plot
sns.regplot(x="Weight_pounds", y="Price", data=df)
plt.ylim(0,)

(0.0, 3974.15)
```



```
# Correlation values of the three attributes with Price
for param in ["CPU_frequency", "Screen_Size_inch", "Weight_pounds"]:
    print(f"Correlation of Price and {param} is ",
          df[[param, "Price"]].corr())
```


```
Correlation of Price and CPU_frequency is
CPU_frequency    Price
CPU_frequency    1.000000  0.366666
Price            0.366666  1.000000
Correlation of Price and Screen_Size_inch is
Screen_Size_inch    Price
Screen_Size_inch    1.000000 -0.110644
Price              -0.110644  1.000000
Correlation of Price and Weight_pounds is
Weight_pounds    Price
Weight_pounds    1.000000 -0.050312
Price           -0.050312  1.000000
```

Interpretation: "CPU\_frequency" has a 36% positive correlation with the price of the laptops. The other two parameters have weak correlation with price.


## Categorical features

Generate Box plots for the different feature that hold categorical values. These features would be "Category", "GPU", "OS", "CPU\_core", "RAM\_GB", "Storage\_GB\_SSD"

```
# Write your code below and press Shift+Enter to execute
# Category Box plot
sns.boxplot(x="Category", y="Price", data=df)
<AxesSubplot:xlabel='Category', ylabel='Price'>
```




```
# Write your code below and press Shift+Enter to execute
# GPU Box plot
sns.boxplot(x="GPU", y="Price", data=df)
<AxesSubplot:xlabel='GPU', ylabel='Price'>
```




```
# Write your code below and press Shift+Enter to execute
# OS Box plot
sns.boxplot(x="OS", y="Price", data=df)
<AxesSubplot:xlabel='OS', ylabel='Price'>
```






```
# Write your code below and press Shift+Enter to execute
# CPU_core Box plot
sns.boxplot(x="CPU_core", y="Price", data=df)

<AxesSubplot:xlabel='CPU_core', ylabel='Price'>
```




```
# Write your code below and press Shift+Enter to execute
# RAM_GB Box plot
sns.boxplot(x="RAM_GB", y="Price", data=df)
<AxesSubplot:xlabel='RAM_GB', ylabel='Price'>
```



```
# Write your code below and press Shift+Enter to execute
# Storage_GB_SSD Box plot
sns.boxplot(x="Storage_GB_SSD", y="Price", data=df)

<AxesSubplot:xlabel='Storage_GB_SSD', ylabel='Price'>
```



## Task 2 - Descriptive Statistical Analysis

Generate the statistical description of all the features being used in the data set. Include "object" data types as well.

```
# Write your code below and press Shift+Enter to execute
print(df.describe())
print(df.describe(include=['object']))
```

	Unnamed: 0.1	Unnamed: 0	Category	GPU	OS \
count	238.000000	238.000000	238.000000	238.000000	238.000000
mean	118.500000	118.500000	3.205882	2.151261	1.058824
std	68.848868	68.848868	0.776533	0.638282	0.235790
min	0.000000	0.000000	1.000000	1.000000	1.000000
25%	59.250000	59.250000	3.000000	2.000000	1.000000
50%	118.500000	118.500000	3.000000	2.000000	1.000000
75%	177.750000	177.750000	4.000000	3.000000	1.000000
max	237.000000	237.000000	5.000000	3.000000	2.000000

	CPU_core	Screen_Size_inch	CPU_frequency	RAM_GB \
count	238.000000	238.000000	238.000000	238.000000
mean	5.630252	14.688655	0.813822	7.882353
std	1.241787	1.166045	0.141860	2.482603

min	3.000000	12.000000	0.413793	4.000000
25%	5.000000	14.000000	0.689655	8.000000
50%	5.000000	15.000000	0.862069	8.000000
75%	7.000000	15.600000	0.931034	8.000000
max	7.000000	17.300000	1.000000	16.000000

	Storage_GB_SSD	Weight_pounds	Price	Screen-Full_HD \
count	238.000000	238.000000	238.000000	238.000000
mean	245.781513	4.106221	1462.344538	0.676471
std	34.765316	1.078442	574.607699	0.468809
min	128.000000	1.786050	527.000000	0.000000
25%	256.000000	3.246863	1066.500000	0.000000
50%	256.000000	4.106221	1333.000000	1.000000
75%	256.000000	4.851000	1777.000000	1.000000
max	256.000000	7.938000	3810.000000	1.000000

	Screen-IPS_panel
count	238.000000
mean	0.323529
std	0.468809
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

	Manufacturer	Price-binned
count	238	238
unique	11	3
top	Dell	Low
freq	71	160

## Task 3 - GroupBy and Pivot Tables

Group the parameters "GPU", "CPU\_core" and "Price" to make a pivot table and visualize this connection using the pcolor plot.

```
# Write your code below and press Shift+Enter to execute
# Create the group
df_gptest = df[['GPU', 'CPU_core', 'Price']]
grouped_test1 =
df_gptest.groupby(['GPU', 'CPU_core'], as_index=False).mean()
print(grouped_test1)
```

	GPU	CPU_core	Price
0	1	3	769.250000
1	1	5	998.500000
2	1	7	1167.941176

3	2	3	785.076923
4	2	5	1462.197674
5	2	7	1744.621622
6	3	3	784.000000
7	3	5	1220.680000
8	3	7	1945.097561

*# Write your code below and press Shift+Enter to execute*

*# Create the Pivot table*

```
grouped_pivot = grouped_test1.pivot(index='GPU',columns='CPU_core')
print(grouped_pivot)
```

	Price		
CPU_core	3	5	7
GPU			
1	769.250000	998.500000	1167.941176
2	785.076923	1462.197674	1744.621622
3	784.000000	1220.680000	1945.097561

*# Write your code below and press Shift+Enter to execute*

*# Create the Plot*

```
fig, ax = plt.subplots()
im = ax.pcolor(grouped_pivot, cmap='RdBu')
```

*#label names*

```
row_labels = grouped_pivot.columns.levels[1]
col_labels = grouped_pivot.index
```

*#move ticks and labels to the center*


```
ax.set_xticks(np.arange(grouped_pivot.shape[1]) + 0.5, minor=False)
ax.set_yticks(np.arange(grouped_pivot.shape[0]) + 0.5, minor=False)
```

*#insert labels*

```
ax.set_xticklabels(row_labels, minor=False)
ax.set_yticklabels(col_labels, minor=False)
```

```
fig.colorbar(im)
```

```
<matplotlib.colorbar.Colorbar at 0xa2147e0>
```



## Task 4 - Pearson Correlation and p-values

Use the `scipy.stats.pearsonr()` function to evaluate the Pearson Coefficient and the p-values for each parameter tested above. This will help you determine the parameters most likely to have a strong effect on the price of the laptops.

```
# Write your code below and press Shift+Enter to execute
for param in
['RAM_GB', 'CPU_frequency', 'Storage_GB_SSD', 'Screen_Size_inch', 'Weight_
pounds', 'CPU_core', 'OS', 'GPU', 'Category']:
    pearson_coef, p_value = stats.pearsonr(df[param], df['Price'])
    print(param)
    print("The Pearson Correlation Coefficient for ", param, " is",
pearson_coef, " with a P-value of P =", p_value)
```

```
RAM_GB
The Pearson Correlation Coefficient for RAM_GB is 0.5492972971857849
with a P-value of P = 3.6815606288424503e-20
CPU_frequency
The Pearson Correlation Coefficient for CPU_frequency is
0.3666655589258861 with a P-value of P = 5.50246335071342e-09
Storage_GB_SSD
The Pearson Correlation Coefficient for Storage_GB_SSD is
0.24342075521810297 with a P-value of P = 0.00014898923191724168
```

```
Screen_Size_inch
The Pearson Correlation Coefficient for Screen_Size_inch is -
0.11064420817118291 with a P-value of P = 0.08853397846830661
Weight_pounds
The Pearson Correlation Coefficient for Weight_pounds is -
0.050312258377515455 with a P-value of P = 0.4397693853433894
CPU_core
The Pearson Correlation Coefficient for CPU_core is
0.45939777733551174 with a P-value of P = 7.912950127008979e-14
OS
The Pearson Correlation Coefficient for OS is -0.22172980114827356
with a P-value of P = 0.0005696642559246817
GPU
The Pearson Correlation Coefficient for GPU is 0.2882981988881427
with a P-value of P = 6.166949698364507e-06
Category
The Pearson Correlation Coefficient for Category is
0.286242755812641 with a P-value of P = 7.225696235806858e-06
```

# Congratulations! You have completed the lab

## Authors

[Abhishek Gagneja](#)

[Vicky Kuo](#)

## Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2023-09-15	0.1	Abhishek Gagneja	Initial Version Created
2023-09-18	0.2	Vicky Kuo	Reviewed and Revised

Copyright © 2023 IBM Corporation. All rights reserved.